

Classification-based reinforcement learning

Vincent Antaki

McGill University

Classification-based reinforcement learning

Approach based on :

- ▶ Lagoudakis & Parr, *Reinforcement Learning as Classification : Leveraging Modern Classifiers*
- ▶ Riedmiller, *Neural Fitted Q Iteration - First Experiences with a Data Efficient Neural Reinforcement Learning Method*
- ▶ Farahmand, Precup, Barreto, Ghavamzadeh
Classification-based Approximate Policy Iteration : Experiments and Extended Discussions
- ▶ Lagoudakis & Parr, *Least-Squares Policy Iteration*

Classification-based reinforcement learning

Why classification-based reinforcement learning?

- ▶ Attempt to leverage advantages of supervised learning for reinforcement learning problems (ex. data efficiency, handling of non-linearity)
- ▶ Find structure directly in the action space

Policy iteration

```
>>> While not satisfied with the policy :  
...     policy_eval()  
...     policy_update()
```

Figure: The underlying principles behind policy iteration

Algorithm CAPI(Π, ν, K)

Input: Policy space Π , State distribution ν , Number of iterations K

Initialize: Let $\pi_{(0)} \in \Pi$ be an arbitrary policy

for $k = 0, 1, \dots, K - 1$ **do**

Construct a dataset $\mathcal{D}_n^{(k)} = \{X_i\}_{i=1}^n, X_i \stackrel{\text{i.i.d.}}{\sim} \mathbb{K}$

$\hat{Q}^{\pi_k} \leftarrow \text{PolicyEval}(\pi_k)$

$\pi_{k+1} \leftarrow \arg\min_{\pi \in \Pi} \hat{L}_n^{\pi_k}(\pi)$ (action-gap-weighted classification)

end for

Figure: The CAPI framework

* Image from Farahmand, Precup, Barreto, Ghavamzadeh *Classification-based Approximate Policy Iteration : Experiments and Extended Discussions*

Action gap :

- ▶ Given state X_i and an action a , we consider the absolute difference between $\hat{Q}^{\pi_*}(X_i, a)$ and $\hat{Q}^{\pi_*}(X_i, a^*)$
- ▶ When action gap is low, regret for choosing the non-optimal action is low and confusion is more likely to happen.

Very important :

$$\hat{L}_n^{\pi_k}(\pi) = \sum_{X_i \in \mathcal{D}_n^{(k)}} \mathbf{g}_{\hat{Q}^{\pi_k}}(X_i) \mathbb{I}\{\pi(X_i) \neq \operatorname{argmax}_{a \in \mathcal{A}} \hat{Q}^{\pi_k}(X_i, a)\}$$

Policy evaluation

We want to approximate π^* by learning Q -values from our samples. Our options include :

- ▶ Rollout
- ▶ Least-square Temporal Difference Q-learning
- ▶ Neural Fitted Q iteration

Rollout

```
Rollout ( $\mathcal{M}, s, a, \gamma, \pi, K, T$ )  
  //  $\mathcal{M}$  : Generative model  
  //  $(s, a)$  : State-action pair whose value is sought  
  //  $\gamma$  : Discount factor  
  //  $\pi$  : Policy  
  //  $K$  : Number of trajectories  
  //  $T$  : Length of each trajectory  
  
  for  $k = 1$  to  $K$   
     $(s', r) \leftarrow \text{SIMULATE}(\mathcal{M}, s, a)$   
     $\tilde{Q}_k \leftarrow r$   
     $s \leftarrow s'$   
    for  $t = 1$  to  $T - 1$   
       $(s', r) \leftarrow \text{SIMULATE}(\mathcal{M}, s, \pi(s))$   
       $\tilde{Q}_k \leftarrow \tilde{Q}_k + \gamma^t r$   
       $s \leftarrow s'$   
  
   $\tilde{Q} \leftarrow \frac{1}{K} \sum_{k=1}^K \tilde{Q}_k$   
  
  return  $\tilde{Q}$ 
```

Figure: The rollout algorithm

► Simulation-based approximation.

LSTDQ

```
LSTDQ ( $D, k, \phi, \gamma, \pi$ )           // Learns  $\hat{Q}^\pi$  from samples

//  $D$  : Source of samples  $(s, a, r, s')$ 
//  $k$  : Number of basis functions
//  $\phi$  : Basis functions
//  $\gamma$  : Discount factor
//  $\pi$  : Policy whose value function is sought

 $\tilde{\mathbf{A}} \leftarrow \mathbf{0}$            //  $(k \times k)$  matrix
 $\tilde{\mathbf{b}} \leftarrow \mathbf{0}$        //  $(k \times 1)$  vector

for each  $(s, a, r, s') \in D$ 
     $\tilde{\mathbf{A}} \leftarrow \tilde{\mathbf{A}} + \phi(s, a) \left( \phi(s, a) - \gamma \phi(s', \pi(s')) \right)^\top$ 
     $\tilde{\mathbf{b}} \leftarrow \tilde{\mathbf{b}} + \phi(s, a) r$ 

 $\tilde{\mathbf{w}}^\pi \leftarrow \tilde{\mathbf{A}}^{-1} \tilde{\mathbf{b}}$ 

return  $\tilde{\mathbf{w}}^\pi$ 
```

Figure: The LSTDQ algorithm

- ▶ Linear function approximation.
- ▶ Need to use basis functions.
- ▶ Requires pseudo-matrix inversions.

```

NFQ_main() {
  input: a set of transition samples  $D$ ; output: Q-value function  $Q_N$ 
  k=0
  init_MLP()  $\rightarrow Q_0$ ;
  Do {
    generate_pattern_set  $P = \{(input^l, target^l), l = 1, \dots, \#D\}$  where:
       $input^l = s^l, u^l$ ,
       $target^l = c(s^l, u^l, s^l) + \gamma \min_b Q_k(s^l, b)$ 
    Rprop_training( $P$ )  $\rightarrow Q_{k+1}$ 
    k:= k+1
  } WHILE ( $k < N$ )

```

- ▶ Use a neural net for regression
- ▶ Trained with SGD or an RProp variant.

* Image from Riedmiller, *Neural Fitted Q Iteration - First Experiences with a Data Efficient Neural Reinforcement Learning Method*

Policy update

Using our approximation of π^* and our samples, we generate multiple examples from each seen state.

- ▶ We give label 0 to for pairs (s, a^*)
- ▶ We give label 1 to for pairs $(s, a) \forall a \neq a^*$

Then, we train a classifier on the dataset, set a tie-breaker policy and we have a new policy.

Technical consideration

We should technically end our algorithm when the policy converges or when the preset maximum number of iterations is reached.

- ▶ Threshold on empirical similarity between policies as stopping criterion.

Advantages

What are the possible advantages to use a classifier for the policy update?

- ▶ Data-efficient methods.
- ▶ Lots of options to handle non-linearity.
- ▶ Can detect structure inherent to the action space.

Advantages

What are the possible advantages to use a classifier for the policy update?

- ▶ Data-efficient methods.
- ▶ Lots of options to handle non-linearity.
- ▶ Can detect structure inherent to the action space.
- ▶ We can make our implementation Scikit compatible.

Load balancing problem 1

- ▶ The agent has 4 'servers' at its disposition and needs to dispatch them tasks it receives. Tasks arrive randomly following a Poisson distribution with $\lambda = 2$.
- ▶ Tasks requires a certain amount of work to be completed. This amount of work required to complete a task is equal to $1 + T$ where T is a random poisson variable with $\lambda = 4$. The agent never knows the associated workload with a task.
- ▶ All server queues have a maximum length of 10. If the agent tries to add a task to an already full queue, the task is discarded and the agents receive a -50 points reward.
- ▶ At every timestep, all servers accomplish one unit of work on the first task in their queue.
- ▶ Upon the completion a task by a server, the agent receives a reward equal to $\frac{5}{\text{\# of iteration to complete task}}$.

Load balancing problem 2

- ▶ The amount of work generated by the servers is now different for every server and stochastic.
- ▶ Distributions : $\mathcal{N}(0.9, 0.1), \mathcal{N}(1, 0.1), \mathcal{N}(1.1, 0.1), \mathcal{N}(1, 0.25)$

Load balancing problem 3

- ▶ Every server has a "heat index" which is between 0 and a certain upper bound.
- ▶ The heat index increases with a fixed rate for every timestep the server is working.
- ▶ The heat index decreases with a (higher) fixed rate for every timestep the server is not working.
- ▶ The amount of work generated by the servers is reduced proportionally to the heat index down to a minimum of 80% of its original capacity.
- ▶ The agent must learn to give short break to servers if possible.

Considerations

For simplicity, we consider the state to be defined as follows :

- ▶ Current timestep
- ▶ Queue length of all servers
- ▶ Time since the current task has been added to the queue for all servers
- ▶ Number of timesteps spent on current task for all servers.

Considerations

Difficulties inherent to that problem :

- ▶ Delayed reward
- ▶ A bit of stochasticity
- ▶ Partially observable state
- ▶ Average reward problem?

Considerations

The optimal policy is however very simple.

- ▶ Give the task to the server which is expected to complete it the soonest.

For the first variant, this means :

- ▶ Give the task to the server with the shortest queue.
- ▶ If multiples servers have the shortest length of queue, give the task to the one which been running his task for the longest amount of time

Methodology and other technical considerations

- ▶ We terminate simulation after a maximum of 500 decisions.
- ▶ We generate batches of 50 episodes for iteration of the main loop (with $\epsilon = 0.2$).
- ▶ Examples are discarded after each policy update.
- ▶ We monitor the reward for each episode and the number of time the agent tries to add a task to an already filled queue.
- ▶ We use a onehot encoding of the action when learning the classifier

Baselines

In term of baselines, we have :

- ▶ The random agent
- ▶ The optimal agent for rproblem variant 1
- ▶ LSPI
- ▶ NFQ

Approaches

Previous approach used :

- ▶ Hand-designed features
- ▶ LSPI

Current approach uses :

- ▶ NFQ with 2 layers MLP with relu activation, no activation on the output neuron, L1 and L2 regularisation.
- ▶ Another MLP for policy update.

Approaches

Possibles add-on to our approach :

- ▶ Shared structure between both MLP
- ▶ Experience replay

The End

Thank you!